

Generative AI 12-Week Learning Roadmap (Beginner to Intermediate)

Week-by-Week Index:

1. **Week 1: Introduction to Generative AI** – Overview of generative AI concepts, examples, and applications.
2. **Week 2: Machine Learning & Deep Learning Foundations** – Essential ML/DL background for understanding generative models.
3. **Week 3: Text Generation Fundamentals (NLP Basics)** – How language models work (from n-grams & RNNs to the idea of Transformers).
4. **Week 4: Introduction to Transformers and Language Models** – Understanding Transformers, attention mechanism, and modern LLMs.
5. **Week 5: Autoencoders and Variational Autoencoders (VAEs)** – Learning latent representations and basic generative modeling for images.
6. **Week 6: Generative Adversarial Networks (GANs)** – GAN architecture, training process, and image generation applications.
7. **Week 7: Diffusion Models and Advanced Image Generation** – Modern diffusion-based image generators (e.g. Stable Diffusion) and how they work.
8. **Week 8: Prompt Engineering and LLM Applications** – Developing skills to interact with and harness large language models effectively.
9. **Week 9: Building Generative AI Projects (Hugging Face & Gradio)** – Hands-on creation of simple generative AI apps and demos.
10. **Week 10: Fine-Tuning Generative Models** – Customizing pre-trained models (LLMs or image models) on new data (e.g. via fine-tuning, LoRA).
11. **Week 11: Deployment and Scaling of Generative AI** – Deploying generative models in real-world applications, optimization and inference considerations.
12. **Week 12: Ethics, Bias, and Future Trends in Generative AI** – Responsible AI usage, ethical concerns, societal impact, and emerging trends.

Week 1: Introduction to Generative AI

Study Goals: By the end of Week 1, you should understand what *generative AI* is and why it's important. You will learn how generative AI differs from other AI (discriminative models), see exciting examples (e.g. ChatGPT, DALL·E), and become familiar with key domains (text, images, audio). The aim is to get a high-level overview and motivation for deeper study.

Topics Covered:

- Definition of Generative AI and how it *creates* new content (text, images, etc.) rather than just analyzing data ¹ ² .
- Major *use cases* of generative models in 2025: chatbots, image generation, music, code generation, etc.
- Real-world impact and trends – how generative AI is transforming industries (media, education, business, etc.) ³ ⁴ .

- Examples of popular generative AI tools (ChatGPT, Stable Diffusion, etc.) to illustrate capabilities ⁵
⁶ .

Resources:

- **Video:** *Intro to Generative AI – MIT CMS.590 (2023 Lecture)* – A beginner-friendly introduction explaining what generative AI is and showcasing examples of AI-generated content (text and visuals). This lecture sets the stage for why generative models matter and what they can do (free on YouTube, ~30 min).
- **Course:** *Generative AI for Everyone (DeepLearning.AI, 2023)* – An introductory course by Andrew Ng covering how generative AI works and what it can and can't do ⁷ . It provides an overview of modern generative AI tools with real-world examples and hands-on exercises in using them ⁸ . *Format:* ~3 hours, non-technical, free to audit. Great for getting a broad understanding of generative AI's capabilities in work and life.
- **Article:** *"Generative AI: The Ultimate Beginner's Guide (2025)" – OdinSchool* – A comprehensive beginner's article that defines generative AI in simple terms and discusses its rise in 2025 ¹ . It explains differences between AI, machine learning, and generative AI, and highlights popular models like GPT-4 and DALL·E 3 ⁹ ¹⁰ . (*Reading tip:* Focus on sections "What is Generative AI?" and real-world applications.)
- **Video:** *"Two Minute Papers – OpenAI GPT-3: Good at Almost Everything!"* – A short, engaging video (~5 min) that demonstrates the power of a generative language model (GPT-3) with fun examples. This will give you a sense of the *remarkable outputs* generative AI can produce in text form, serving as motivation for the coming weeks.

Week 2: Machine Learning & Deep Learning Foundations

Study Goals: Build the necessary foundation in machine learning (ML) and deep learning to understand how generative models function. By week's end, you should grasp key concepts like neural networks, how models learn from data, and fundamentals of training (without going too deep into math). This background will make the mechanics of generative models (in later weeks) much clearer.

Topics Covered:

- Basics of machine learning: training vs. inference, datasets (training/validation), and the concept of a model learning patterns from data.
- Introduction to neural networks: neurons, layers, weights – how a simple network takes inputs and produces outputs by successive transformations.
- Key concepts in deep learning: activation functions, loss function and optimization (gradient descent/backpropagation) to adjust weights.
- Overfitting and generalization basics (why models need the right capacity and regularization). Understanding that generative models are a subset of ML that *learn distributions* of data.

Resources:

- **Video:** *3Blue1Brown – "But what is a Neural Network?"* – A renowned visual introduction to neural networks (YouTube, 20 min). Grant Sanderson illustrates how a basic neural network recognizes

handwriting, building intuition for how networks adjust weights to learn ¹¹. This helps demystify the “magic” behind AI by showing step-by-step how networks learn.

- **Video:** *freeCodeCamp – “Deep Learning Crash Course”* – A beginner-friendly crash course (YouTube, ~2 hours) that efficiently introduces the most important deep learning concepts ¹². It covers how neural networks are trained, different architectures, and basic terminology. *Tip:* You don’t need to absorb every detail; focus on big-picture concepts like layers, training loops, and what makes deep learning powerful.
 - **Course (Optional):** *Neural Networks and Deep Learning (Coursera by Andrew Ng)* – If you prefer a structured course, this is Course 1 of the Deep Learning Specialization. It’s more detailed, spanning several weeks of content on neural network basics and optimization. You can audit for free. Consider reviewing Week 1–2 of this course for a deeper dive into how neural networks learn (if time permits).
 - **Article/Guide:** *“Neural Networks” – Chapter from fast.ai book or online course* – A practical explanation of neural network fundamentals with code examples. Fast.ai’s material is known for being beginner-friendly yet practical. Skim a chapter or lesson that introduces how to train a simple network (e.g., recognizing MNIST digits) to reinforce concepts through an example.
-

Week 3: Text Generation Fundamentals (NLP Basics)

Study Goals: This week, you will focus on *natural language generation*. The goal is to learn how early language models worked and set the stage for modern approaches. By the end, you should understand the concept of a language model (predicting the next word/sequence) and know why sequence data is special. You’ll also get an introduction to recurrent neural networks (RNNs) and see their limitations, which will hint at why Transformers were needed.

Topics Covered:

- What is a language model? Understanding how generative text models predict the next word or character based on previous context. Simple examples like **n-gram** models leading into neural approaches.
- Recurrent Neural Networks (RNNs): how they process sequences by maintaining a hidden state. Concepts of sequence dependence, and an intro to **LSTM** (Long Short-Term Memory) networks to handle longer dependencies.
- The challenge of long-term dependencies and problems like vanishing gradients in RNNs ¹³ ¹⁴. Why we needed new architectures for very long texts.
- (Preview) Attention mechanism idea: a teaser that will lead into Transformers next week – understanding that models can “attend” to important bits of input rather than just sequentially reading.

Resources:

- **Lecture:** *MIT 6.S191 (2025) – Recurrent Neural Networks and Transformers* – The first part of this MIT lecture (~58 min total) covers RNN fundamentals ¹³. It explains sequence modeling, unfolding an RNN through time, and demonstrates an RNN predicting text ¹⁵. You’ll also learn about LSTM networks and see examples of RNN applications. (*Note:* The latter part of the lecture introduces attention and Transformers – which is perfect to watch as a segue into Week 4) ¹⁴ ¹⁶.

- **Video:** *StatQuest – “Variational Autoencoders EXPLAINED”* – Wait, this is mis-filed for text generation. (We will replace this with an NLP resource.)
- **Video:** *StatQuest – “RNNs and LSTMs”* – A gentle, visual explanation (YouTube, ~15 min) of how RNNs handle sequences and what LSTMs add. Josh Starmer (StatQuest) breaks down the gating mechanisms of LSTMs in simple terms. This will reinforce the MIT lecture content with a different explanation style, ensuring you grasp why handling long-term context in text is hard.
- **Hands-on Tutorial:** *“Text Generation with Python (Char-RNN)”* – A step-by-step tutorial (blog or Jupyter Notebook) where you train a simple character-level RNN to generate text (e.g., Shakespeare-like text). Following such a tutorial will solidify how sequence models are trained. For instance, **Karpathy’s “min-char-rnn”** is a classic resource (if available in a simplified form) or consider a Hugging Face Transformers quick tour using an RNN model. *Goal:* Get a feel for training and sampling from a language model, even a very small one.
- **Fun Video:** *Two Minute Papers – “Training an AI to Write Shakespeare”* – (Hypothetical example) A short video showing results from a language model that learned to imitate Shakespeare. While not an actual Two Minute Papers title, seek out a brief video that shows an *entertaining result* of text generation. It will illustrate the progress from nonsense output to coherent style mimicry, keeping you motivated.

(Note: Focus on understanding concepts over coding this week, but a little experimentation can be enlightening if you have programming background.)

Week 4: Introduction to Transformers and Language Models

Study Goals: This week you dive into the revolution in NLP: the *Transformer* architecture. By the end, you should understand the key ideas behind Transformers (self-attention, parallel processing) and why they outperform RNNs for language modeling. You’ll also learn about large language models (LLMs) like GPT, BERT (though BERT is not generative, it’s good context), and how Transformers enabled the current generative AI boom. The goal is an intermediate understanding of **how modern LLMs work** at a conceptual level.

Topics Covered:

- Sequence modeling with **attention**: The self-attention mechanism that lets models weigh the relevance of different words in a sequence (the famous “Attention Is All You Need” concept).
- The Transformer architecture: idea of an encoder-decoder vs. decoder-only (GPT) models. Multi-head attention, positional embeddings, and why Transformers can handle long context better than RNNs by design.
- Overview of notable Transformer-based models: GPT series for text generation, BERT for understanding (mention for contrast), and newer hybrid models. How GPT-3 and GPT-4 are essentially very large Transformers pretrained on internet-scale data.
- Emergence of capabilities with scale: how increasing model size and data led to surprising abilities (few-shot learning, etc.). Also touch on the concept of fine-tuning vs. prompting these large models, which leads into later weeks.

Resources:

- **Lecture:** MIT 6.S191 – “Attention and Transformers” (2025 Edition) – The second half of the MIT lecture from last week (Lecture 2) or a dedicated lecture on Transformers. It explains the intuition of attention and shows how a Transformer processes a sequence in parallel rather than step-by-step ¹⁴ ¹⁶. Key points include understanding *why* attention can capture long-range dependencies and a quick look at how Transformers are applied in language models (this lecture is ~58 min; focus on the latter sections covering attention and Transformer architecture).
 - **Article:** “The Illustrated Transformer” by Jay Alammar – A highly-recommended blog post that visually breaks down the Transformer model. It walks through an example of translating a sentence, illustrating how self-attention works and how the model builds representations. This resource is great to solidify your conceptual model of attention. (Expect diagrams of word relationships, etc., making it easier to digest the math-heavy paper in an intuitive way.)
 - **YouTube Playlist:** “Transformers from Scratch” (Jan 2023) – A short series (if available) or single video that builds a toy Transformer model. Some creators (like AI coffee-break style channels) have 10–15 min explainers on how Transformers encode positions and mix information. Use this to reinforce the *mechanics* of Transformers in simpler terms after reading the Alammar article.
 - **Interactive Tutorial:** Hugging Face Course – Chapter on “Attention and Transformers” – The Hugging Face free course has a section explaining Transformers and even lets you play with code to see how attention weights can be visualized. Completing this chapter (or a similar Colab) will give a practical sense of what happens inside a Transformer when processing text.
 - **Resource:** Two Minute Papers – “OpenAI’s GPT-4: Absolutely Amazing!” – A light, enthusiastic overview of the capabilities of a large Transformer-based model (GPT-4) from Two Minute Papers (YouTube, ~10 min). It doesn’t teach architecture, but it will show *outcomes* – such as GPT-4 solving tasks – which reinforces why learning about Transformers is worth it. Seeing what GPT-4 can do (e.g., code, answer complex questions) will contextualize the theory with real-world impact.
-

Week 5: Autoencoders and Variational Autoencoders (VAEs)

Study Goals: Shift focus to *generative modeling in vision*. This week covers **Autoencoders** and **Variational Autoencoders**, which are fundamental generative models for images (and other data). By week’s end, you should understand how autoencoders compress data into a latent representation, and how VAEs can generate new data by sampling from a latent space. These concepts build intuition for latent spaces and will also be useful later (even in understanding diffusion models and GANs).

Topics Covered:

- Autoencoders (AE): Encoder-decoder neural networks that compress input data into a smaller *latent vector* and then reconstruct the input. Use cases: dimensionality reduction, denoising. Key point: they learn a latent representation but typically *do not* generate new data (they just reconstruct).
- **Variational Autoencoders (VAE):** Extending autoencoders to generative models by learning a *distribution* over the latent space ¹⁷ ¹⁸. The encoder produces a mean and variance (a distribution of latent vectors) instead of a single point. By sampling from this latent distribution and decoding, VAEs can generate new, similar data ¹⁹.
- VAE specifics: the **reparameterization trick** (to allow gradient flow despite sampling) – conceptually

understand it introduces randomness without breaking backpropagation ²⁰. Loss function combining reconstruction error and KL-divergence (to keep latent distribution close to a normal prior).

- Limitations of VAEs: Generated outputs tend to be blurry for images (since optimizing for average outcome), but they are stable to train and good for understanding latent spaces. Mention how VAEs paved the way for richer latent generative models and are still used in e.g. diffusion model pipelines as part of image decoders.

Resources:

- **Video:** *"Autoencoders Explained"* – Valerio Velardo (*The Sound of AI*) – A beginner-friendly video (~15 min) explaining the concept of autoencoders. It uses analogies and simple visuals to show how an image (or sound) can be compressed to a smaller vector and then reconstructed. This gives an intuition for the encoder-decoder structure and the idea of a *bottleneck* layer capturing the essence of the data.
- **Video:** *StatQuest – "Variational Autoencoders (VAEs) Explained"* – A clear explanation (YouTube, ~17 min) where StatQuest walks through what makes a VAE different from a standard autoencoder. It covers the idea of sampling from a latent space and how VAEs can generate new outputs by interpolating between known data points. By the end of this, terms like "latent space" and "variational inference" will feel much less intimidating.
- **Lecture:** *MIT 6.S191 (2025) – Deep Generative Modeling (VAE segment)* – In MIT's Intro to DL course Lecture 4, the first part is devoted to VAEs ²¹. The lecturer explains latent variable models and how VAEs generate new samples resembling the training distribution ²¹. Watching this 15–20 minute segment reinforces the concepts with an academic perspective, including how VAEs were a milestone in generative modeling.
- **Article:** *IBM Developer – "What is a Variational Autoencoder?"* – A written tutorial/guide that defines VAEs in accessible language. It explains how VAEs encode a **continuous, probabilistic latent space** that enables generating new data ¹⁷ ¹⁸. It also touches on applications of VAEs (anomaly detection, etc.) and provides nice conceptual summaries (e.g., the bridge weight analogy for latent variables) ²² ²³. Use this for reference if any concept from the videos is unclear, and to see formal definitions of terms.
- **Mini-Project:** *Train a VAE on MNIST (Optional)* – If you're inclined and have some coding skills, use a PyTorch or TensorFlow tutorial (many exist) to train a VAE on the MNIST digits dataset. It's a classic exercise where you can visualize the latent space (perhaps 2D) and see the decoder generate digit images from sampled latent points. This hands-on experience will make the theoretical aspects concrete – you'll see, for example, how interpolation in latent space produces smooth morphing between digits, illustrating the VAE's generative power.

Week 6: Generative Adversarial Networks (GANs)

Study Goals: This week you'll tackle **GANs**, one of the most famous generative model frameworks, especially for images. By the end of the week, you should understand the two-player game between a Generator and a Discriminator in GANs, how GANs can produce remarkably realistic images, and the challenges in training them. You'll also explore improvements over the basic GAN and see some state-of-the-art results (which will be both inspiring and a precursor to diffusion models).

Topics Covered:

- GAN basics: the concept of two networks contesting – **Generator** produces fake data (e.g. images), **Discriminator** tries to distinguish fakes from reals. Training is a minimax game where the generator tries to fool the discriminator ²¹. When successful, the generator outputs look like real data.
- Loss formulation: understanding at a high-level the adversarial loss (generator trying to minimize discriminator success, discriminator trying to maximize it). No need for equations, but grasp that it's different from traditional losses and can be unstable.
- Training difficulties: mode collapse (generator output diversity falls), non-convergence, and tricks that help (feature matching, one-sided label smoothing, etc. – just mention a couple as examples). Emphasize that training GANs is an art due to the delicate balance.
- Notable GAN architectures: **DCGAN** (deep convolutional GAN – first to generate good faces with conv nets), **Pix2Pix/CycleGAN** (conditional GANs for image-to-image translation), and **StyleGAN** (state-of-the-art by NVIDIA for photorealistic faces). These show the evolution and potency of GANs.
- Example outputs: “This person does not exist” (AI-generated faces) and other eye-catching results of GANs to appreciate the quality and to understand why GANs were a breakthrough in generative imaging.

Resources:

- **Course:** *Generative Adversarial Networks (GANs) Specialization* – *DeepLearning.AI* – A three-course Coursera specialization (audit for free) that provides an **exciting introduction to image generation with GANs** ²⁴. In Course 1, you learn GAN fundamentals and build a basic GAN; Course 2 covers advanced topics (evaluating GAN outputs, GAN hacks, and introduces StyleGAN techniques) ²⁵ ²⁶; Course 3 explores GAN applications (like image translation with Pix2Pix/CycleGAN) and ethical considerations (bias in GAN outputs, privacy) ²⁷ ²⁸. *You may not complete all courses now*, but reviewing the Course 1 content this week (lectures on GAN basics by Sharon Zhou) will give you a solid grounding. The hands-on assignments use PyTorch to implement GAN components, which is great practical learning if you have the time.
- **Video:** *“GANs Explained with Image Examples”* – (YouTube, e.g. Yannic Kilcher or similar) – A 10–15 minute explainer video that visually walks through the GAN mechanism. It often uses a simple analogy (e.g., art forger (generator) vs. detective (discriminator)) to explain the adversarial training loop. This will reinforce conceptually *how* GANs learn to generate data that mimics a target distribution.
- **Lecture:** *MIT 6.S191 – Deep Generative Modeling (GAN segment)* – In the MIT Lecture 4 (Deep Generative Modeling), after VAEs the instructor discusses GANs ²¹. This segment (~20 min) covers how GANs approach the generative task differently (no explicit latent distribution enforcement like VAE, instead a game between networks). It's a concise academic explanation of GAN architecture and key issues, solidifying your theoretical understanding.
- **Article:** *“GANs in 50 lines of code (PyTorch)”* – *Medium or Blog* – A walkthrough where you implement a simple GAN (e.g., generating MNIST digits). This resource shows code but focuses on explaining each part (generator network, discriminator network, training loop alternating between them). Following this will demystify the training process. You'll see how the loss drives the generator to improve. Even if you don't code it yourself, stepping through the logic is valuable to connect theory with practice.
- **Video (Fun):** *Two Minute Papers – “None of These Faces Are Real!”* – A short video showcasing results from NVIDIA's StyleGAN (which generates incredibly realistic human faces) ²⁹. It's an inspiring look at what GANs can achieve. The video briefly explains that these faces are AI-

generated and touches on the idea of a latent space that the GAN is sampling to produce novel images. Watching this will both amaze you and cement why GANs have been a hot topic – they produced a leap in output quality (you can literally see it).

Week 7: Diffusion Models and Advanced Image Generation

Study Goals: This week focuses on *diffusion models*, the cutting-edge approach behind tools like **Stable Diffusion** and **DALL-E 3** for image generation. By the end, you should understand the basic idea of diffusion (gradually noising images then learning to reverse that process), why this approach produces high-quality results, and how it differs from GANs. You'll also get introduced to using diffusion models (possibly via open-source implementations) and appreciate their capabilities in generating not just images but other modalities (with the concept being extendable).

Topics Covered:

- Diffusion model concept: the two processes – a **forward process** that adds noise to data step by step until it becomes pure noise, and a **reverse process** where a model is trained to gradually remove noise, reconstructing data. Essentially, the model learns to denoise noisy images in steps, which allows it to generate images from noise.
- Key elements: **U-Net** architecture often used as the denoiser, **schedule of noise** addition/removal, and the notion of *sampling* (creating new data by starting with random noise and applying the learned denoising steps).
- Why diffusion models are powerful: They tend to produce more stable and diverse results than GANs (mode coverage, less training instability) at the cost of being slower to sample (iterative process). Also mention that they scale well and can incorporate guidance (like text prompts via techniques like CLIP guidance or classifier-free guidance).
- **Stable Diffusion** as a case study: an open-source image diffusion model that takes a text prompt and generates images. Explain at a high level how it uses a text encoder (from a model like CLIP) to influence the image generation.
- Other applications: Briefly note that diffusion isn't just for images – there are diffusion models for audio (e.g., noise to music) and even video (though video diffusion is heavy). But the focus should remain on image generation since it's the most mature use case.

Resources:

- **Course:** “How Diffusion Models Work” (DeepLearning.AI short course) – An intermediate-level short course (approx. 2 hours) taught by Sharon Zhou that will **take you through the steps to build and train a diffusion model** ³⁰. It covers the intuition and math behind diffusion processes, including how to noise images and train a model to reverse it. By following this course, you'll gain a deep familiarity with the diffusion approach and even implement a basic diffusion model yourself. (This is one of the best ways to grasp diffusion – highly recommended if you have the time to code through it.)
- **Article:** AI Summer – “Diffusion Models: The Math from Scratch” – A written guide that, as the title suggests, walks through diffusion model formulation step-by-step. It starts from the basics of adding Gaussian noise and then explains the training objective for denoising. While it includes math, it also provides intuition in each section. Use this to supplement the course if you want clarity on the underlying equations and to solidify your understanding of terms like *variational lower-bound*, *score matching*, etc., in a more approachable way.

- **Tutorial:** *Hugging Face – “Introduction to Diffusers” (Unit 1)* – Hugging Face’s Diffusers library is a high-level API to use diffusion models. This introductory unit shows **how diffusion models work and how to create your own using Diffusers** ³¹. It likely includes a hands-on example with code where you start from noise and generate an image, using a pre-trained model for parts of the pipeline. Running through this notebook will give you practical experience with diffusion in code and familiarize you with using pre-trained models like Stable Diffusion in a few lines of code.
 - **Video:** *Two Minute Papers – “Stable Diffusion: DALL-E 2 for Free, for Everyone!”* – A short video (~7 min) that introduced Stable Diffusion when it came out. It explains how this model brought image generation to the masses (being open-source and runnable on consumer GPUs) and shows example outputs. This will help you understand the significance of diffusion models in the context of generative AI history, and seeing the impressive images will reinforce what the technology can do (the video underscores that Stable Diffusion can create images from text prompts, similar to DALL-E, but *open and free* ³²).
 - **Project:** *Play with Stable Diffusion* – Use an accessible tool to try generating images yourself. For instance, try **Stable Diffusion via Hugging Face Spaces** (there are web demos) or install the Diffusers library to run a prompt or two (if you have the hardware). Even using a cloud service or a Colab notebook for Stable Diffusion is an option. The goal is to connect theory with real output: write a few fun prompts and see the images that come out. As you do this, consider how the model is applying what you learned – it’s running the reverse diffusion process conditioned on your text prompt. This hands-on play will make the learning concrete and delight you with the creative possibilities.
-

Week 8: Prompt Engineering and LLM Applications

Study Goals: Now that you have learned how generative models work under the hood, Week 8 shifts to a *practical skill*: using generative AI models effectively via **prompt engineering**. By the end of this week, you should know how to craft prompts to get better outputs from language models, understand concepts like zero-shot vs. few-shot prompting, and be aware of the capabilities and limitations of LLMs in various applications. You’ll also explore some popular interfaces and frameworks for working with LLMs (like ChatGPT, OpenAI API, etc.). This week is about becoming a savvy *user* of generative AI, which complements your technical understanding.

Topics Covered:

- **Prompt Engineering:** Techniques to communicate with LLMs. This includes choosing the right instructions, providing context or examples (few-shot learning), and formatting outputs (for example, asking for answers in a structured format). Realize that the prompt is essentially programming the model in natural language.
- Examples of effective prompts vs. poor prompts: You’ll see case studies where a small change in phrasing yields a significantly better answer from the AI. Introduction to strategies like role prompting (“Act as a tutor...”), constraint prompting (“Answer in JSON format...”), etc.
- Notion of model *temperature* and randomness in outputs – understanding how you can tweak generation settings (if using an API) to influence creativity vs. consistency.
- Overview of LLM-powered applications: chatbots, writing assistants, code generation (like GitHub Copilot), etc., and how prompt engineering is critical in each. Possibly touch on prompt chaining or using frameworks (LangChain) to manage multi-step prompts and incorporate tools, but focus on core prompting first.

- Responsible use in prompting: guidelines like avoiding sensitive or biased prompt content if you want safe outputs, and the idea of the models' guardrails (they might refuse certain prompts, etc.). This segues gently into ethics next weeks.

Resources:

- **Course:** *"ChatGPT Prompt Engineering for Developers"* (DeepLearning.AI, 2023) – A free short course by OpenAI and Andrew Ng that teaches prompt engineering best practices. It covers how to write prompts for different tasks (summarization, brainstorming, coding help, etc.), and includes hands-on examples with the OpenAI API. This course will give you practical techniques (like how adding few-shot examples can guide the model) and also provides insight into how to handle when the model refuses or gives wrong answers. It's concise (~1.5 hours) and very relevant 33
- **Guide:** *OpenAI Cookbook – "Best practices for prompt design"* – OpenAI's documentation/guides often have a section on how to phrase instructions for models like GPT-3/4. This resource outlines tips such as being explicit, giving system/user role directives (if using ChatGPT format), providing examples in the prompt, etc. Reading through these best practices will reinforce what you learned in the course above and serve as a reference for future prompting tasks.
- **YouTube:** *"Prompt Engineering 101"* – (e.g., a talk by Isa Fulford or another OpenAI expert) – Look for a YouTube video around 20–30 minutes that is a talk or tutorial on prompt engineering. Many have been produced given the interest in ChatGPT. A good one will show interactive demos: e.g., taking a simple prompt that produces a mediocre result, then iteratively refining it into a great prompt. This live process helps you learn *how to think* when crafting prompts.
- **Article:** *"Awesome Prompt Engineering Techniques" (collection)* – A curated list or blog post that enumerates different prompt patterns. For instance, it might list techniques like **Chain-of-Thought prompting** (instruct the model to reason step by step) and give examples, or **Self-Consistency** prompting (have the model generate multiple solutions and pick the best). Skimming such a list can introduce you to advanced methods that you might try out. It's not necessary to master all of them now, but knowing what's possible is useful as you continue to experiment with LLMs.
- **Practice Tool:** *OpenAI Playground or Hugging Face Chat Interface* – Spend some time this week practicing prompts with a real LLM. You can use ChatGPT itself or an open-source model on a platform like Hugging Face Chat. Try tasks like: *writing a short story in a specific style, summarizing a technical article, answering a question with a chain-of-thought*, etc. Apply what you learned: adjust the prompt, add examples, etc., and see how the output changes. This kind of hands-on practice is perhaps the most effective way to solidify prompt engineering skills.

Week 9: Building Generative AI Projects (Hugging Face & Gradio)

Study Goals: In Week 9, you will learn how to *apply* generative AI by building simple projects and demos. The focus is on using pre-trained models (from Hugging Face or similar) and creating an interface, rather than training models from scratch. By the end of the week, you should be comfortable finding models for a task (like text generation, image generation, etc.), using libraries (Hugging Face Transformers or Diffusers) to invoke those models in a few lines of code, and wrapping them in a simple web app interface (using **Gradio** or Streamlit). Essentially, you'll go from "I know how this model works"

to “I can make something with it that others can interact with.” This is an important step toward deployment (which is next week).

Topics Covered:

- Hugging Face Hub: how to search for and download pre-trained generative models. Understanding model cards, trying out tasks with the **Pipeline API** (e.g., `pipeline("text-generation", model="gpt2")` or `pipeline("image-generation", model="stability-ai/stable-diffusion")`).
- Using models programmatically: loading a model and tokenizer in Python, feeding input and getting output. Maybe a quick example of generating text with GPT-2 or images with Stable Diffusion (with appropriate safety and small scale).
- **Gradio** introduction: a Python library that makes it trivial to create a web UI for any ML model. You define a function (input -> output) and Gradio helps serve it with a textbox, image box, etc. The goal is to learn how to take your generative model function and expose it so that non-technical users (or you yourself) can interact with it easily.
- Hugging Face Spaces: a free hosting for Gradio apps. Understanding that you can push your Gradio demo to Spaces (a git repo) and it will be live on the web. This ties into deployment next week, but here it's more about prototyping quickly.
- Example projects could include: a simple chatbot using an open-source smaller LLM, an image generation app where user enters a prompt and sees an AI-generated image, an audio generation demo, etc. The emphasis is on *breadth* – see that with these tools, you can make a lot of AI ideas real in a short time.

Resources:

- **Short Course:** “Open Source Models with Hugging Face” (DeepLearning.AI) – This free course (about 2 hours) is perfect for this week. It teaches you how to **find and filter models on Hugging Face Hub, use the transformers library to perform NLP, audio, image, and multimodal tasks with just a few lines of code, and easily share your models via Gradio and Hugging Face Spaces** ³³ ³⁵. It's a step-by-step guide through multiple tasks: turning a language model into a chatbot, doing translation, image captioning, even combining modalities (like describing an image) ³⁶ ³⁷. Crucially, it has you deploy an app on Spaces by the end ³⁸. By following this course, you will essentially hit all the week's goals with hands-on examples.
- **Documentation:** *Hugging Face Transformers official docs* – Specifically the Quickstart section. Read the parts about using pipelines for text generation and maybe the chapter on “Deployment” if available. The docs will reinforce how to load a model and run it. It's also a good reference for later when you experiment with different models.
- **Gradio Tutorial:** “Gradio + HuggingFace Spaces: A Tutorial” by Tanishq Abraham – A blog post that shows how to create a Gradio app and deploy on Spaces in <10 minutes ³⁹. It covers creating a new Space, selecting Gradio, adding your code, and seeing it live. This tutorial demystifies the process and will be useful when you try to share your own mini-project. The PyImageSearch tutorial on deploying Gradio apps ⁴⁰ is another alternative with more step-by-step screenshots if needed.
- **Example Repository:** *Hugging Face Spaces – “Hello World” generative app* – Find a simple open-source Space (on huggingface.co/spaces) that implements a basic generative task. For instance, there might be a Space that does “text to emoji art” or a basic story generator. Look at the code (usually just a `app.py` or `notebook.ipynb`) to see how they structure the Gradio interface

and call the model. Learning from examples will give you ideas on how to structure your own projects.

- **Project Ideas:** If you have time, try to build one small project yourself. Some ideas:
 - A **text completion app**: Using GPT-2 or another model to finish a sentence or paragraph the user starts.
 - An **image generation app**: Using Stable Diffusion (if you have the resources) or a smaller model to generate an image from a prompt.
 - A **poetry generator**: Fine-tune (or prompt) a model to produce haikus or rap lyrics.
 - Use Gradio to make it interactive and deploy on Spaces. Even if it's simple, you'll learn a ton by doing – and it's rewarding to have a shareable link to your AI creation!

(By building something this week, you'll be well-prepared for the final steps of fine-tuning and full deployment, having already gotten your hands dirty with real models and apps.)

Week 10: Fine-Tuning Generative Models

Study Goals: Having used pre-trained models as-is, this week you'll learn how to *fine-tune* generative models on custom data. By week's end, you should understand why fine-tuning is useful (adapting a model to a specific style or domain), the general process for fine-tuning both an LLM and an image generative model, and the tools/techniques that make fine-tuning feasible (like using smaller adapters or low-rank adaptation for huge models). You'll ideally run or walkthrough one end-to-end example of fine-tuning (e.g., fine-tune a small language model on a niche text dataset, or fine-tune Stable Diffusion on a specific art style or on your own images). This will solidify how you can customize generative AI beyond prompting alone.

Topics Covered:

- Fine-tuning concept: taking a pre-trained model and further training it on a smaller dataset for a specific task or style. How it differs from training from scratch (needs much less data and time because the model already learned a lot).
- Fine-tuning LLMs: methods like full fine-tuning (updating all model weights, feasible for smaller models), and **parameter-efficient tuning** like **LoRA (Low-Rank Adaptation)** or prompt tuning for very large models (which inject a small trainable layer or embeddings instead of updating the whole model)
41 42 . Why techniques like LoRA are popular – they make fine-tuning large models possible on consumer hardware by training only a fraction of parameters.
- Example fine-tuning scenario: e.g., fine-tune a GPT-2 model to generate Shakespearean-style text, or fine-tune a smaller open-source LLM on your own conversation data to create a domain-specific chatbot. Steps involve preparing the dataset, using a library (Hugging Face Trainer or LoRA libraries), training for a few epochs, and then testing the adapted model.
- Fine-tuning image models: introduction to **DreamBooth** (fine-tuning Stable Diffusion to output images of a new subject, like your face or a specific character) and other fine-tuning use-cases (e.g., training a Stable Diffusion model to adopt a particular art style by providing examples). Mention that image model fine-tuning often requires less data (tens of images) but careful training to avoid overfitting, and often uses techniques like attention control or adding new embeddings (like Textual Inversion for new concepts).
- Caution: Overfitting and evaluation – how to tell if your fine-tuned model is actually good or just memorized your training data (especially for text, check if it's just regurgitating). Also, the concept of *catastrophic forgetting* – if not careful, a fine-tuned model might lose some of its general ability (less of an issue with methods like LoRA which preserve the original weights).

Resources:

- **Course Module:** *“Finetuning Large Language Models”* – A short course or module (perhaps from Hugging Face or DeepLearning.AI) that covers the full fine-tuning workflow ⁴³. It should discuss selecting a base model, preparing data, training, and deploying the fine-tuned model. For example, Hugging Face’s course might have a chapter on fine-tuning causal language models for text generation, and DeepLearning.AI’s Generative AI series might include a segment on fine-tuning LLMs with tools like LlamaIndex or others. This will give you a structured understanding and possibly a unified example to follow.
- **Blog Tutorial:** *“Fine-tuning GPT-2 for Text Generation”* – A how-to blog (e.g., on Medium or Hugging Face Blog) that walks through fine-tuning GPT-2 on a custom text dataset. It typically covers: using Hugging Face Transformers Trainer, setting up your dataset (perhaps a collection of poems or dialogues), training the model, and then generating samples from the fine-tuned model to see the new style. Following such a tutorial (or at least reading it) will clarify the practical steps and common pitfalls (like needing to adjust the learning rate to avoid the model just outputting gibberish or copying training data). ⁴⁴ (Hugging Face’s official documentation on fine-tuning is a great reference here ⁴⁴).
- **Video:** *“Low-Rank Adaptation (LoRA) Explained”* – A concise video or talk (10 min) that explains how LoRA fine-tunes large models by injecting small trainable weight matrices. This will help you conceptually understand how we can fine-tune a 7 billion parameter model by training maybe just 10 million parameters inserted into it. Understanding LoRA and similar techniques is key to working with models like LLaMA or GPT-J in resource-constrained environments ⁴¹. The video might also mention quantization (reducing model precision to make it smaller) ⁴¹, which is another helpful concept for deployment.
- **Tutorial:** *“Fine-tune Stable Diffusion with DreamBooth”* – Many guides exist (on Hugging Face blog or others) that show how to use DreamBooth. DreamBooth typically involves providing ~5–10 images of a subject and fine-tuning the Stable Diffusion model so that subject can be inserted into any generated scene via a special token. This tutorial will give you insight into fine-tuning an image generation model. It usually uses a Colab notebook due to the GPU requirement. Even if you don’t run it, read through to understand the process (it’s surprisingly straightforward: load model, add your images as dataset, train for a few hundred steps, use the new model). It’s a fun and illustrative example of customizing a powerful generative model.
- **Community Forums/Documentation:** *Hugging Face Discuss – Fine-tuning tips* – Skimming through a forum thread or FAQ on fine-tuning issues (for example, someone asking “My fine-tuned model just outputs the same thing over and over – what did I do wrong?”) can prepare you for common challenges. It exposes you to troubleshooting mindset: maybe the learning rate was too high, or dataset too small. This will deepen your understanding and ensure you don’t treat fine-tuning as a black box.

(By the end of this week, try to have at least one fine-tuned outcome to show, even if small – for instance, a GPT-2 fine-tuned to output a specific style. It’s a rewarding experience and consolidates a lot of what you learned about both modeling and application.)

Week 11: Deployment and Scaling of Generative AI

Study Goals: In Week 11, you'll learn how to *deploy* generative AI models so that they can be used in real-world applications reliably. This involves understanding the infrastructure side: choosing where to host models, how to serve them (API endpoints, etc.), and considerations for scaling (latency, throughput). By the end of the week, you should know the options for deploying a model (cloud services, Hugging Face Inference API, on-device), understand how to optimize models for faster inference (like quantization, distillation), and be aware of monitoring and maintaining a deployed model (since generative models can have unique issues like sometimes producing inappropriate content that you need to filter). Essentially, this week bridges the gap from a prototype (like the Gradio demo you made) to a production-ready service.

Topics Covered:

- Deployment basics: wrapping your model in an API. For example, using FastAPI or Flask to create a web service that accepts requests (like a text prompt) and returns generated output. Or using cloud functions / specialized ML serving platforms.
- Utilizing cloud ML services: Many cloud providers (AWS, GCP, Azure) have services for deploying models. For instance, Amazon SageMaker or Google's Vertex AI – they handle provisioning GPU instances and exposing endpoints. You should learn at a high level how one might containerize a model and deploy on such a service. (You might not do it hands-on due to cost, but understanding is key.)
- **Hugging Face Spaces** vs production APIs: HF Spaces is great for demos but not ideal for heavy production use or custom scaling. For production, one might use Hugging Face Inference Endpoints (a managed solution) or host a model on their own server. Learn the differences in reliability, scalability, and cost.
- Model optimization for deployment: Introduce techniques like **quantization** (reducing precision of model weights from 32-bit to 8-bit, etc., to make inference faster) ⁴¹ ⁴², **knowledge distillation** (compressing a large model into a smaller one by training a “student” model to imitate the “teacher”), and leveraging specialized hardware (GPUs, TPUs, or new AI chips). The idea isn't to master these, but to know they exist and why they help (e.g., quantized models can run with less memory and sometimes little loss in quality).
- Scaling and monitoring: If 1000 users hit your model at once, how do you handle it? This involves scaling out (multiple instances, load balancing). Also, monitoring usage and performance (did latency spike? did the model output any disallowed content?). Covering a bit about logging and content filtering (perhaps using OpenAI's moderation model or similar if you deploy a text model) as part of a responsible deployment.

Resources:

- **Guide:** *Hugging Face – “Deploying models” (Official docs or blog)* – Hugging Face provides guides on deploying models, either via their Inference API or on your own. A specific guide might be “How to deploy a Transformer model to production”. It would walk you through packaging the model (maybe using Docker), setting up an API, and pinging it with requests. This is a concrete resource that can serve as a reference blueprint for deployment.
- **Video/Workshop:** *“Building and Deploying LLMs at Scale”* – Look for a recorded workshop or conference talk (~30-60 min) where practitioners discuss deploying large models. For example, a talk where an engineer from OpenAI or Cohere discusses how they serve thousands of requests, or a tutorial from a cloud provider on serving stable diffusion in a scalable way. This will give you real-world context and tips (like using batching of requests to better utilize GPU, etc.). It might

also mention tools like Ray Serve, MLflow, or Kubernetes for deploying ML models, giving you pointers for further exploration.

- **Short Course:** *“Efficiently Serving LLMs”* – A course or piece of one (like Pedram’s Pedestal or others) focusing on serving large models efficiently ⁴¹. From the guide we saw, this covers techniques like LoRA and quantization in a deployment context. It may also mention frameworks optimized for inference (like ONNX Runtime or TensorRT for optimizing model graphs). Studying this will make you aware of how to reduce inference time and memory usage – crucial for deploying generative models which are often huge and slow.
- **Tutorial:** *“Deploy Gradio app on Hugging Face Spaces”* – While you did this in Week 9, here consider it as a stepping stone to more robust deployment. The tutorial by PyImageSearch ⁴⁰ or the official Gradio docs (Sharing Your App) ⁴⁵ show how to deploy an app. Revisit your Week 9 project: could it handle more users? Maybe read about how to attach a GPU to a Space or use a paid tier if needed. The idea is to think: what would it take to go from a fun demo to a service people rely on? Often that means moving to a more controlled environment, but this reflection is valuable.
- **Case Study:** *Blog post “Scaling Stable Diffusion for millions of users”* – Hypothetical, but there might be blog posts from Stability AI or others on how they served Stable Diffusion via API. For instance, learn how Stability’s DreamStudio or NightCafe backend might work. These case studies give practical insight (e.g., running multiple model instances, using CPU offloading for diffusion steps, etc.).
- **Checklist:** *“Production Readiness for Generative Models”* – Find or create a checklist of things to consider (some engineering blogs or Medium posts by companies have this). It might include: security (ensure the model can’t be easily reverse-engineered or doesn’t expose something sensitive), updating models (how to deploy a new model version with minimal downtime), setting up monitoring alerts if something goes wrong, etc. It’s a nice way to summarize the deployment mindset.

(This week is fairly broad and might feel more like engineering/devops than pure ML, but it’s crucial for turning your knowledge into real services. Don’t worry if you can’t practice everything – focus on understanding the landscape and knowing where to look when you do need to deploy.)

Week 12: Ethics, Bias, and Future Trends in Generative AI

Study Goals: The final week zooms out to consider the *bigger picture* of generative AI. By the end, you should be aware of the major ethical issues (bias, misinformation, copyright, etc.) associated with generative models, and understand the importance of responsible AI practices. You’ll also take stock of current trends and what’s on the horizon: e.g., even larger multimodal models, better alignment techniques, and how society and policy might evolve around generative AI. This week isn’t about technical skills but about developing a critical perspective and foresight, which is essential for any AI practitioner or informed user.

Topics Covered:

- **Bias in generative models:** How models can perpetuate or amplify societal biases present in training data. Examples: image generators often reflecting gender/racial stereotypes in outputs ⁴⁶ ⁴⁷ (e.g., when asked for images of a CEO vs. a nurse), or language models producing biased or toxic text.

Understand that generative AI learns from vast data that includes biased content, and without intervention, will reflect that ⁴⁸ ⁴⁹ .

- **Hallucinations and misinformation:** Especially with LLMs – the tendency to produce confident-sounding but incorrect information ⁵⁰ ⁵¹ . Why this happens (the model's goal is to be plausible, not correct ⁵²) and the risks it poses if people trust AI output blindly. Also, deepfakes and generated images/audio that can spread misinformation, and the need for detection and watermarking techniques.

- **Intellectual Property and Ownership:** Discussion on who owns AI-generated content. If a model was trained on a dataset of artists' work or on copyrighted text, is its output derivative? Current debates about copyright lawsuits (e.g., authors vs. OpenAI, or artists vs. Stable Diffusion) and how companies are responding (some models limiting training data, offering opt-outs, etc.).

- **Privacy:** Generative models like LLMs might regurgitate parts of training data, which could include personal data. And when fine-tuning on company data, one must ensure no sensitive info leaks out in generations. Acknowledge techniques like prompt filtering and red-teaming that try to catch these issues.

- **Responsible deployment:** Connecting with Week 11 – measures like content filters (OpenAI's moderation API, Stability's safety classifier for images) to prevent misuse. The importance of human-in-the-loop for critical applications (e.g., AI-generated content should be reviewed in medical or legal contexts).

- **AI alignment and future directions:** Mention efforts to make AI follow human ethical guidelines (RLHF – Reinforcement Learning from Human Feedback – used in ChatGPT's training to curb bad outputs). And new research like constitutional AI (Anthropic) or tool use (retrieval-augmented generation to reduce hallucinations).

- **Future trends:** Larger and *multimodal models* (like GPT-4 vision, or Google's Gemini that combines language and images). Also, more efficient models (distilled or specialized models on edge devices). And the role of regulation: EU AI Act, FTC interest in generative AI – how the landscape might get rules to ensure safety and fairness.

- End on a note that generative AI is a powerful tool that must be used responsibly – as a newly educated practitioner, you should carry forward an ethical mindset in your projects.

Resources:

- **Article:** *MIT Sloan – “When AI Gets It Wrong: Addressing AI Hallucinations and Bias”* – This article provides an overview of pitfalls in generative AI outputs, particularly biased and inaccurate content ⁴⁶ ⁴⁷ . It gives real examples (e.g., Stable Diffusion amplifying stereotypes ⁴⁷ , ChatGPT producing racist outputs in a study ⁵³) and discusses why these issues occur ⁵¹ ⁴⁹ . It also offers strategies to mitigate these problems ⁵⁴ ⁵⁵ , like critically evaluating AI outputs and using diverse sources. This is an excellent summary to ground your understanding of ethical challenges and how we might navigate them.

- **Report:** *“GPT-4 System Card” (OpenAI)* – OpenAI published a system card for GPT-4 outlining the steps they took to make it safer, including the types of problematic content it can produce and how they mitigated it. Skim this (it's a bit long) to see a concrete example of an ethics review: it covers misuse risks (like generating harmful advice), bias evaluations, and the model's limitations. This will show you the thoroughness that leading organizations are applying to deploy models responsibly.

- **Course/Module:** *“AI Ethics and Society” (short course or Coursera module)* – If time permits, doing a short AI ethics course or at least the sections relevant to generative AI is valuable. For example, Elements of AI (a free course) has an ethics chapter. Or Coursera's “AI For Everyone” by Andrew Ng has a week on ethical issues. Focus on topics like fairness, accountability, transparency in AI

decisions, and apply those concepts to generative scenarios (like the need to disclose AI-generated content to maintain transparency).

- **Panel Discussion (Video): “Ethics of Generative AI”** – Many conferences have panel talks with experts (AI researchers, ethicists, artists, etc.) discussing the implications of generative AI. Find a recent one on YouTube (for instance, a SXSW or CogX panel from 2023). This will expose you to multiple viewpoints – technical, social, legal. Hearing experts talk about, say, deepfakes in politics or AI in education (students using ChatGPT) will enrich your perspective beyond the technical realm.
- **Future Trend Blog: “The Future of Generative AI”** – A forward-looking piece by an AI luminary (could be a blog by Andrej Karpathy, or an IEEE Spectrum article). It might speculate on where this is going: integration of modalities (text+image+audio models), more personalized models (everyone having their own AI), and necessary breakthroughs (like solving factual accuracy, or reducing data hunger through better training methods). Reading such an article in the final week is inspiring and will help you consolidate what you’ve learned in context – understanding not just what generative AI is today, but where it’s heading and how you can be part of that future.

Congratulations! Over these 12 weeks, you’ve progressed from a beginner to an intermediate level in Generative AI. You started with foundational concepts and gradually built up to advanced techniques and practical skills. You’ve learned not only how generative models work (from RNNs and Transformers to GANs and Diffusion models) but also how to use them, adapt them, deploy them, and critically evaluate their outputs.

As you move forward, remember to keep learning incrementally – the field is evolving quickly. Continue experimenting with new models on Hugging Face, follow AI news (Two Minute Papers will keep you updated on cool new research!), and perhaps dive deeper into specialized areas you found interesting (like music generation or model interpretability). And always keep ethics in mind: generative AI will significantly impact society, so use your skills responsibly and for positive purposes.

Good luck on your generative AI journey! The roadmap you’ve completed is a strong foundation, and there’s so much more you can build on it.

1 2 3 4 5 6 9 10 **Generative AI: The Ultimate Beginner’s Guide (2025 Edition) | Odinschool**
<https://www.odinschool.com/blog/generative-ai-the-ultimate-beginners-guide-2025-edition>

7 8 **Generative AI for Everyone - DeepLearning.AI**
<https://www.deeplearning.ai/courses/generative-ai-for-everyone/>

11 12 **Mastering Generative AI: A Roadmap from Zero to Expertise in Gen AI field | by Incletech Admin | Medium**
<https://medium.com/@incle/mastering-generative-ai-a-roadmap-from-zero-to-expertise-in-gen-ai-field-95a058defcda>

13 14 15 16 **Free Video: Recurrent Neural Networks and Transformers from Alexander Amini | Class Central**
<https://www.classcentral.com/course/youtube-mit-6-s191-recurrent-neural-networks-and-transformers-128080>

17 18 19 20 22 23 **What is a Variational Autoencoder? | IBM**
<https://www.ibm.com/think/topics/variational-autoencoder>

21 Free Video: Deep Generative Modeling - MIT 6.S191 Lecture 4 from Alexander Amini | Class Central
<https://www.classcentral.com/course/youtube-mit-6-s191-deep-generative-modeling-438815>

24 25 26 27 28 Generative Adversarial Networks (GANs) | Coursera
<https://www.coursera.org/specializations/generative-adversarial-networks-gans>

29 None of These Faces Are Real! - YouTube
<https://www.youtube.com/watch?v=-cOYwZ2XcAc>

30 How Diffusion Models Work - DeepLearning.AI
<https://www.deeplearning.ai/short-courses/how-diffusion-models-work/>

31 Unit 1: An Introduction to Diffusion Models - Hugging Face
<https://huggingface.co/learn/diffusion-course/en/unit1/1>

32 Stable Diffusion: DALL-E 2 For Free, For Everyone! - YouTube
<https://www.youtube.com/watch?v=nVhmFski3vg>

33 34 35 36 37 38 Open Source Models with Hugging Face - DeepLearning.AI
<https://www.deeplearning.ai/short-courses/open-source-models-hugging-face/>

39 Gradio + HuggingFace Spaces: A Tutorial – Dr. Tanishq Abraham
https://www.tanishq.ai/blog/gradio_hf_spaces_tutorial/

40 Deploy Gradio Apps on Hugging Face Spaces - PyImageSearch
<https://pyimagesearch.com/2024/12/30/deploy-gradio-apps-on-hugging-face-spaces/>

41 42 43 Your Guide to Generative AI Courses - DeepLearning.AI
<https://www.deeplearning.ai/resources/generative-ai-courses-guide/>

44 Fine-tuning - Hugging Face
<https://huggingface.co/docs/transformers/en/training>

45 Sharing Your App - Gradio
<https://www.gradio.app/guides/sharing-your-app>

46 47 48 49 50 51 52 53 54 55 When AI Gets It Wrong: Addressing AI Hallucinations and Bias - MIT Sloan Teaching & Learning Technologies
<https://mitsloanedtech.mit.edu/ai/basics/addressing-ai-hallucinations-and-bias/>